

DELIVERABLE 1

Jesús García Ayala

Contents

1. Introduction	2
2. Data Font (where is data from?)	4
3. Metadata	5
4. Data description	8
4.1 Dataset dimensions	8
4.2 Types of variables	8
4.3 Descriptive statistics	8
4.4 Preliminary insights	8
5. Preprocessing	27
5.1 Explanation	31
5.1.1 Issues Detected in the Raw Dataset	31
5.1.2 Step-by-Step Preprocessing Explained	33
5.2 Conclusion	35
5.3 Creation of New Variables:	36
5.3.1 Second-Generation Variables:	36
5.3.2 Third-Generation Variables:	38
6. Basic description of post processing in this variables with changes	42
6.1 Main changes in variables	42
6.2 Visualization of changes	43
6.3 Statistics' summary after processing	43
6.4 Conclusions	45

1. Introduction

In today's competitive market, understanding customer behavior is crucial for businesses to optimize their marketing strategies and enhance customer retention. For this reason, we have selected this dataset, as we believe it is highly valuable.

This dataset was chosen because it allows us to analyze key factors influencing customer decision-making, such as household income, family structure, spending patterns, and response to promotional campaigns. By exploring this data, we can identify trends, segment customers, and develop data-driven strategies to improve marketing effectiveness.

Additionally, the dataset contains a variety of numerical and categorical variables, making it an excellent choice for applying preprocessing techniques. This diversity enables us to utilize different analytical methods, including statistical analysis.

By working with this dataset, we aim to gain a deeper understanding of how marketing efforts impact different customer groups, how businesses can enhance their strategies to maximize engagement and revenue, and how to effectively analyze a dataset, including both its preprocessing and postprocessing stages.

Database context:

The database is part of a case study designed to simulate the real challenges faced by data analysts at iFood, the leading food delivery platform in Brazil. The company operates in over a thousand cities and serves millions of customers annually. Maintaining a high level of customer engagement is essential to consolidating its market leadership.

The iFood data team must work on open-ended analytical projects to identify business opportunities, optimize marketing strategies, and support data-driven decision-making. We will put ourselves in the shoes of iFood's data team.

The dataset provides simulated information on customer profiles based on real data and their interactions with iFood's marketing campaigns, including socio-demographic and firmographic data from 2,240 customers, randomly selected for a pilot marketing campaign.

The main challenge is to develop a predictive model that more accurately identifies customers with a higher likelihood of purchasing, allowing iFood to focus its efforts on the most responsive segments.

By applying data-driven strategies, iFood aims to refine its marketing approach, reduce inefficiencies, and strengthen its relationship with customers, solidifying its position as the leader in the food delivery industry.

Objective of the project

The goal of this project is to analyze customer behavior in a sales company by exploring their demographic characteristics, purchasing habits, and response to marketing campaigns. The aim is to extract useful insights that can help improve commercial strategies and better personalize offers for customers.

Methods used

- Exploratory Data Analysis (EDA) to understand data distribution and potential anomalies.
- Handling missing values and outliers to ensure data quality.
- Comparison of variables before and after preprocessing to observe the impact on analysis.
- Use of **R** for visualizations and basic predictive modeling on the dataset.

2. Data Font (where is data from?)

<https://www.kaggle.com/datasets/jackdaoud/marketing-data?resource=download&select=dictionary.png>

The dataset used comes from a customer database of a company that offers various types of products. It contains information about:

- **Customer profile** (income, age, marital status, education level).
- **Purchasing habits** (products bought, purchase channels).
- **Response to marketing campaigns** (acceptance of promotions).
- **Interaction with the service** (website visits, recorded complaints).

The original data has been preprocessed to remove inconsistencies and improve interpretability, as detailed in the following sections.

3. Metadata

This table provides an explanation of all the variables in our dataset, including the variable name, data type, variable type classification, description, and possible values.

Variable name	Variable Type	Description	Possibles values
Id	ID	The ID of the customer	Random number between 3-5 digits
Income	Numerical Continuous	Customer's yearly household income.	> 0
Kidhome	Numerical Discrete	Number of small children in the customer's household.	0-5
Teenhome	Numerical Discrete	Number of teenagers in the customer's household.	0-5
Recency	Numerical Discrete	Number of days since the last purchase.	0-365(aprox)
MntWines	Numerical Continuous	Amount spent on wine in the last 2 years.	≥ 0
MntFruits	Numerical Continuous	Amount spent on fruits in the last 2 years.	≥ 0
MntMeatProducts	Numerical Continuous	Amount spent on meat products in the last 2 years.	≥ 0

MntFishProducts	Numerical Continuous	Amount spent on fish products in the last 2 years.	≥ 0
MntSweetProducts	Numerical Continuous	Amount spent on sweet products in the last 2 years.	≥ 0
MntGoldProds	Numerical Continuous	Amount spent on gold products in the last 2 years.	≥ 0
NumDealsPurchases	Numerical Discrete	Number of purchases made with a discount.	≥ 0
NumWebPurchases	Numerical Discrete	Number of purchases made through the website.	≥ 0
NumCatalogPurchases	Numerical Discrete	Number of purchases made using a catalog.	≥ 0
NumStorePurchases	Numerical Discrete	Number of purchases made directly in a store.	≥ 0
NumWebVisitsMonth	Numerical Discrete	Number of visits to the company's website by the customer in the last month.	≥ 0
AcceptedCmp1	Categorical Binary	Whether the customer accepted the offer in the 1st campaign.	0, 1
AcceptedCmp2	Categorical Binary	Whether the customer accepted the offer in the 2nd campaign.	0, 1

AcceptedCmp3	Categorical Binary	Whether the customer accepted the offer in the 3rd campaign.	0, 1
AcceptedCmp4	Categorical Binary	Whether the customer accepted the offer in the 4th campaign.	0, 1
AcceptedCmp5	Categorical Binary	Whether the customer accepted the offer in the 5th campaign.	0, 1
Complain	Categorical Binary	Whether the customer has filed a complaint in the last 2 years.	0, 1
Z_CostContact	Numerical Discrete	Fixed cost of contacting the customer.	3
Z_Revenue	Numerical Discrete	Fixed revenue value.	11
Response	Categorical Binary	Whether the customer accepted the offer in the last campaign.	0, 1
Year_birth	Numerical Discrete	Customer's year of birth.	1900-2020
Dt_Customer	Numerical Discrete	Date when the customer enrolled .	Date between 1-1-1900/31-12-2020
Marital	Categorical Nominal	Whether the customer is divorced.	Divorced, Married, Single, Together, Widow
Education	Categorical Nominal	The level of studies the consumer has.	2nd cycle, Basic, Graduation, Master PhD

4. Data description

Basic descriptive = basic description of the dataset through descriptive statistics and visualizations.

4.1 Dataset dimensions

The original dataset contains 2240 rows and 29 columns, representing different customer attributes.

4.2 Types of variables

The dataset includes:

- **Numerical (quantitative) variables:** such as income, age, spending on different products, number of purchases through each channel, etc.
- **Categorical (qualitative) variables:** such as marital status, education level and customer classification based on marketing campaign responses.

4.3 Descriptive statistics

Basic statistics for numerical variables include:

- **Mean, median, minimum, maximum and standard deviation.**
- **Histograms to analyze distribution.**
- **Boxplots to identify outliers.**

Categorical variables are summarized using:

- **Frequency tables** to examine category distributions.
- **Bar and pie charts** for a visual representation of common values.

4.4 Preliminary insights

- Initial patterns observed in the data are highlighted.
- Variables with extreme or highly skewed values are discussed.
- Initial hypotheses are formulated for further analysis.

Configuración para mejorar la ejecución

```
knitr::opts_chunk$set(echo = TRUE, warning = FALSE, message = FALSE)
```

Cargar dataset

```
ifood <- read.csv("ifood_no_preprocessed.csv", sep=";", header=TRUE)
```

Mostrar las primeras filas

```
head(ifood)
```

##	ID	Year_Birth	Education	Marital_Status	Income	Kidhome	Teenhome	Dt_Customer
## 1	5524	1957	Graduation	Single	58138	0	0	2012-09-04
## 2	2174	1954	Graduation	Single	46344	1	1	2014-03-08
## 3	4141	1965	Graduation	Together	71613	0	0	2013-08-21
## 4	6182	1984	Graduation	Together	26646	1	0	2014-02-10
## 5	5324	1981	PhD	Married	58293	1	0	2014-01-19
## 6	7446	1967	Master	Together	62513	0	1	2013-09-09
##	Recency	MntWines	MntFruits	MntMeatProducts	MntFishProducts	MntSweetProducts		
## 1	58	635	88	546	172	88		
## 2	38	11	1	6	2	1		
## 3	26	426	49	127	111	21		
## 4	26	11	4	20	10	3		
## 5	94	173	43	118	46	27		
## 6	16	520	42	98	0	42		
##	MntGoldProds	NumDealsPurchases	NumWebPurchases	NumCatalogPurchases				
## 1	88		3	8		10		
## 2	6		2	1		1		
## 3	42		1	8		2		
## 4	5		2	2		0		
## 5	15		5	5		3		
## 6	14		2	6		4		
##	NumStorePurchases	NumWebVisitsMonth	AcceptedCmp3	AcceptedCmp4	AcceptedCmp5			
## 1	4		7	0	0	0		
## 2	2		5	0	0	0		
## 3	10		4	0	0	0		
## 4	4		6	0	0	0		
## 5	6		5	0	0	0		
## 6	10		6	0	0	0		
##	AcceptedCmp1	AcceptedCmp2	Complain	Z_CostContact	Z_Revenue	Response		
## 1	0	0	0	3	11	1		
## 2	0	0	0	3	11	0		
## 3	0	0	0	3	11	0		
## 4	0	0	0	3	11	0		
## 5	0	0	0	3	11	0		
## 6	0	0	0	3	11	0		

Ver estructura del dataset

```
str(ifood)
```

```
## 'data.frame':    2240 obs. of  29 variables:
##  $ ID                : int  5524 2174 4141 6182 5324 7446 965 6177 4855 5899 ...
##  $ Year_Birth         : int  1957 1954 1965 1984 1981 1967 1971 1985 1974 1950 ...
##  $ Education          : chr   "Graduation" "Graduation" "Graduation" "Graduation" ...
##  $ Marital_Status     : chr   "Single" "Single" "Together" "Together" ...
##  $ Income              : int  58138 46344 71613 26646 58293 62513 55635 33454 30351 5648 ...
##  $ Kidhome            : int    0 1 0 1 1 0 0 1 1 1 ...
##  $ Teenhome           : int    0 1 0 0 0 1 1 0 0 1 ...
##  $ Dt_Customer        : chr   "2012-09-04" "2014-03-08" "2013-08-21" "2014-02-10" ...
##  $ Recency            : int    58 38 26 26 94 16 34 32 19 68 ...
##  $ MntWines           : int    635 11 426 11 173 520 235 76 14 28 ...
##  $ MntFruits          : int    88 1 49 4 43 42 65 10 0 0 ...
##  $ MntMeatProducts    : int    546 6 127 20 118 98 164 56 24 6 ...
##  $ MntFishProducts    : int    172 2 111 10 46 0 50 3 3 1 ...
##  $ MntSweetProducts   : int    88 1 21 3 27 42 49 1 3 1 ...
##  $ MntGoldProds       : int    88 6 42 5 15 14 27 23 2 13 ...
##  $ NumDealsPurchases  : int    3 2 1 2 5 2 4 2 1 1 ...
##  $ NumWebPurchases    : int    8 1 8 2 5 6 7 4 3 1 ...
##  $ NumCatalogPurchases: int    10 1 2 0 3 4 3 0 0 0 ...
##  $ NumStorePurchases  : int    4 2 10 4 6 10 7 4 2 0 ...
##  $ NumWebVisitsMonth  : int    7 5 4 6 5 6 6 8 9 20 ...
##  $ AcceptedCmp3       : int    0 0 0 0 0 0 0 0 0 1 ...
##  $ AcceptedCmp4       : int    0 0 0 0 0 0 0 0 0 0 ...
##  $ AcceptedCmp5       : int    0 0 0 0 0 0 0 0 0 0 ...
##  $ AcceptedCmp1       : int    0 0 0 0 0 0 0 0 0 0 ...
##  $ AcceptedCmp2       : int    0 0 0 0 0 0 0 0 0 0 ...
##  $ Complain           : int    0 0 0 0 0 0 0 0 0 0 ...
##  $ Z_CostContact      : int    3 3 3 3 3 3 3 3 3 3 ...
##  $ Z_Revenue          : int    11 11 11 11 11 11 11 11 11 11 ...
##  $ Response           : int    1 0 0 0 0 0 0 0 1 0 ...
```

```
## NumWebPurchases NumCatalogPurchases NumStorePurchases NumWebVisitsMonth
## Min. : 0.000 Min. : 0.000 Min. : 0.00 Min. : 0.000
## 1st Qu.: 2.000 1st Qu.: 0.000 1st Qu.: 3.00 1st Qu.: 3.000
## Median : 4.000 Median : 2.000 Median : 5.00 Median : 6.000
## Mean : 4.085 Mean : 2.662 Mean : 5.79 Mean : 5.317
## 3rd Qu.: 6.000 3rd Qu.: 4.000 3rd Qu.: 8.00 3rd Qu.: 7.000
## Max. :27.000 Max. :28.000 Max. :13.00 Max. :20.000
##
## AcceptedCmp3 AcceptedCmp4 AcceptedCmp5 AcceptedCmp1
## Min. :0.00000 Min. :0.00000 Min. :0.00000 Min. :0.00000
## 1st Qu.:0.00000 1st Qu.:0.00000 1st Qu.:0.00000 1st Qu.:0.00000
## Median :0.00000 Median :0.00000 Median :0.00000 Median :0.00000
## Mean :0.07277 Mean :0.07455 Mean :0.07277 Mean :0.06429
## 3rd Qu.:0.00000 3rd Qu.:0.00000 3rd Qu.:0.00000 3rd Qu.:0.00000
## Max. :1.00000 Max. :1.00000 Max. :1.00000 Max. :1.00000
##
## AcceptedCmp2 Complain Z_CostContact Z_Revenue
## Min. :0.00000 Min. :0.000000 Min. :3 Min. :11
## 1st Qu.:0.00000 1st Qu.:0.000000 1st Qu.:3 1st Qu.:11
## Median :0.00000 Median :0.000000 Median :3 Median :11
## Mean :0.01339 Mean :0.009375 Mean :3 Mean :11
## 3rd Qu.:0.00000 3rd Qu.:0.000000 3rd Qu.:3 3rd Qu.:11
## Max. :1.00000 Max. :1.000000 Max. :3 Max. :11
##
## Response
## Min. :0.0000
## 1st Qu.:0.0000
## Median :0.0000
## Mean :0.1491
## 3rd Qu.:0.0000
## Max. :1.0000
##
```

Resumen Estadístico General

```
# Resumen de todas las variables
summary(ifford)
```

```
## ID Year_Birth Education Marital_Status
## Min. : 0 Min. :1893 Length:2240 Length:2240
## 1st Qu.: 2828 1st Qu.:1959 Class :character Class :character
## Median : 5458 Median :1970 Mode :character Mode :character
## Mean : 5592 Mean :1969
## 3rd Qu.: 8428 3rd Qu.:1977
## Max. :11191 Max. :1996
##
## Income Kidhome Teenhome Dt_Customer
## Min. : 1730 Min. :0.0000 Min. :0.0000 Length:2240
## 1st Qu.: 35303 1st Qu.:0.0000 1st Qu.:0.0000 Class :character
## Median : 51382 Median :0.0000 Median :0.0000 Mode :character
## Mean : 52247 Mean :0.4442 Mean :0.5062
## 3rd Qu.: 68522 3rd Qu.:1.0000 3rd Qu.:1.0000
## Max. :666666 Max. :2.0000 Max. :2.0000
## NA's :24
## Recency MntWines MntFruits MntMeatProducts
## Min. : 0.00 Min. : 0.00 Min. : 0.0 Min. : 0.0
## 1st Qu.:24.00 1st Qu.: 23.75 1st Qu.: 1.0 1st Qu.: 16.0
## Median :49.00 Median : 173.50 Median : 8.0 Median : 67.0
## Mean :49.11 Mean : 303.94 Mean : 26.3 Mean : 166.9
## 3rd Qu.:74.00 3rd Qu.: 504.25 3rd Qu.: 33.0 3rd Qu.: 232.0
## Max. :99.00 Max. :1493.00 Max. :199.0 Max. :1725.0
##
## MntFishProducts MntSweetProducts MntGoldProds NumDealsPurchases
## Min. : 0.00 Min. : 0.00 Min. : 0.00 Min. : 0.000
## 1st Qu.: 3.00 1st Qu.: 1.00 1st Qu.: 9.00 1st Qu.: 1.000
## Median :12.00 Median : 8.00 Median : 24.00 Median : 2.000
## Mean : 37.53 Mean : 27.06 Mean : 44.02 Mean : 2.325
## 3rd Qu.: 50.00 3rd Qu.: 33.00 3rd Qu.: 56.00 3rd Qu.: 3.000
## Max. :259.00 Max. :263.00 Max. :362.00 Max. :15.000
##
```

Análisis de Variables Numéricas

```
# Seleccionar variables numéricas
numericas <- sapply(ifood, is.numeric)
numericas <- names(ifood)[numericas]
```

```
# Histograma y boxplot para cada variable numérica
for (var in numericas) {
  cat("Variable -> ", var, "\n\n")

  # Histograma
  hist(ifood[[var]], main=paste("Histograma de", var), col="skyblue", border="black")

  # Boxplot
  boxplot(ifood[[var]], main=paste("Boxplot de", var), col="orange", horizontal=TRUE)


  # Tabla de frecuencias y resumen estadístico

  # Muestra solo los primeros 20 valores


  print(head(ifood[[var]], 20))
  print(summary(ifood[[var]], 20))
  cat("\n\n")
}
```

```
## Variable -> ID
```

Histograma de ID




Boxplot de ID




```
## [1] 5524 2174 4141 6182 5324 7446 965 6177 4855 5899 1994 387 2125 8180 2569
## [16] 2114 9736 4939 6565 2278
##      Min. 1st Qu.  Median     Mean 3rd Qu.     Max.
##         0   2828    5458    5592    8428   11191
##
##
## Variable -> Year_Birth
```

Histograma de Year_Birth




Boxplot de Year_Birth




```
## [1] 1957 1954 1965 1984 1981 1967 1971 1985 1974 1950 1983 1976 1959 1952 1987
## [16] 1946 1980 1946 1949 1985
##   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##   1893   1959   1970    1969   1977    1996
##
##
## Variable -> Income
```

Histograma de Income




Boxplot de Income




```
## [1] 58138 46344 71613 26646 58293 62513 55635 33454 30351 5648 NA 7500
## [13] 63033 59354 17323 82800 41850 37760 76995 33812
##   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.   NA's
##   1730   35303   51382   52247   68522  666666     24
##
##
## Variable -> Kidhome
```

Histograma de Kidhome




Boxplot de Kidhome




```
## [1] 0 1 0 1 1 0 0 1 1 1 1 0 0 1 0 0 1 0 0 1
##   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
## 0.0000 0.0000 0.0000 0.4442 1.0000 2.0000
##
##
## Variable -> Teenhome
```

Histograma de Teenhome




Boxplot de Teenhome




```
## [1] 0 1 0 0 0 1 1 0 0 1 0 0 0 1 0 0 1 0 1 0
##   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
## 0.0000 0.0000 0.0000 0.5062 1.0000 2.0000
##
##
## Variable -> Recency
```

Histograma de Recency




Boxplot de Recency




```
## [1] 58 38 26 26 94 16 34 32 19 68 11 59 82 53 38 23 51 20 91 86
##   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##    0.00   24.00   49.00   49.11   74.00   99.00
##
##
## Variable -> MntWines
```

Histograma de MntWines




Boxplot de MntWines




```
## [1] 635 11 426 11 173 520 235 76 14 28 5 6 194 233 3
## [16] 1006 53 84 1012 4
##   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##    0.00   23.75   173.50   303.94   504.25  1493.00
##
##
## Variable -> MntFruits
```


Histograma de MntFruits




Boxplot de MntFruits




```
## [1] 88 1 49 4 43 42 65 10 0 0 5 16 61 2 14 22 5 5 80 17
## Min. 1st Qu. Median Mean 3rd Qu. Max.
## 0.0 1.0 8.0 26.3 33.0 199.0
##
##
## Variable -> MntMeatProducts
```

Histograma de MntMeatProducts




Boxplot de MntMeatProducts




```
## [1] 546 6 127 20 118 98 164 56 24 6 6 11 480 53 17 115 19 38 498
## [20] 19
## Min. 1st Qu. Median Mean 3rd Qu. Max.
## 0.0 16.0 67.0 166.9 232.0 1725.0
##
##
## Variable -> MntFishProducts
```

Histograma de MntFishProducts




Boxplot de MntFishProducts




```
## [1] 172  2 111 10 46  0 50  3  3  1  0 11 225  3  6 59  2 150  0
## [20] 30
##   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##   0.00   3.00   12.00   37.53  50.00  259.00
##
##
## Variable -> MntSweetProducts
```

Histograma de MntSweetProducts




Boxplot de MntSweetProducts




```
## [1] 88  1 21  3 27 42 49  1  3  1  2  1 112  5  1 68 13 12 16
## [20] 24
##   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##   0.00   1.00   8.00   27.06  33.00  263.00
##
##
## Variable -> MntGoldProds
```

Histograma de MntGoldProds




Boxplot de MntGoldProds




```
## [1] 88 6 42 5 15 14 27 23 2 13 1 16 30 14 5 45 4 28 176
## [20] 39
##   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##   0.00   9.00   24.00   44.02  56.00  362.00
##
##
## Variable -> NumDealsPurchases
```

Histograma de NumDealsPurchases




Boxplot de NumDealsPurchases




```
## [1] 3 2 1 2 5 2 4 2 1 1 1 1 3 1 1 3 2 2 2
##   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##   0.000   1.000   2.000   2.325   3.000  15.000
##
##
## Variable -> NumWebPurchases
```

Histograma de NumWebPurchases




Boxplot de NumWebPurchases




```
## [1] 8 1 8 2 5 6 7 4 3 1 1 2 3 6 1 7 3 4 11 2
##   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##   0.000  2.000  4.000  4.085  6.000  27.000
##
##
## Variable -> NumCatalogPurchases
```

Histograma de NumCatalogPurchases




Boxplot de NumCatalogPurchases




```
## [1] 10 1 2 0 3 4 3 0 0 0 0 0 4 1 0 6 0 1 4 1
##   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##   0.000  0.000  2.000  2.662  4.000  28.000
##
##
## Variable -> NumStorePurchases
```

Histograma de NumStorePurchases




Boxplot de NumStorePurchases




```
## [1] 4 2 10 4 6 10 7 4 2 0 2 3 8 5 3 12 3 6 9 3
##   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##   0.00   3.00   5.00   5.79   8.00   13.00
##
##
## Variable -> NumWebVisitsMonth
```

Histograma de NumWebVisitsMonth




Boxplot de NumWebVisitsMonth




```
## [1] 7 5 4 6 5 6 6 8 9 20 7 8 2 6 8 3 8 7 5 6
##   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##   0.000   3.000   6.000   5.317   7.000   20.000
##
##
## Variable -> AcceptedCmp3
```

Histograma de AcceptedCmp3




Boxplot de AcceptedCmp3




```
## [1] 0 0 0 0 0 0 0 0 0 0 1 0 0 0 0 0 0 0 0 0
##   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
## 0.00000 0.00000 0.00000 0.07277 0.00000 1.00000
##
##
## Variable -> AcceptedCmp4
```

Histograma de AcceptedCmp4




Boxplot de AcceptedCmp4




```
## [1] 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
##   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
## 0.00000 0.00000 0.00000 0.07455 0.00000 1.00000
##
##
## Variable -> AcceptedCmp5
```

Histograma de AcceptedCmp5




Boxplot de AcceptedCmp5




```
## [1] 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 1 0 0 0 0
##   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
## 0.00000 0.00000 0.00000 0.07277 0.00000 1.00000
##
##
## Variable -> AcceptedCmp1
```

Histograma de AcceptedCmp1




Boxplot de AcceptedCmp1




```
## [1] 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 1 0 0 1 0
##   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
## 0.00000 0.00000 0.00000 0.06429 0.00000 1.00000
##
##
## Variable -> AcceptedCmp2
```

Histograma de AcceptedCmp2




Boxplot de AcceptedCmp2




```
## [1] 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
##   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
## 0.00000 0.00000 0.00000 0.01339 0.00000 1.00000
##
##
## Variable -> Complain
```

Histograma de Complain




Boxplot de Complain




```
## [1] 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
##   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
## 0.000000 0.000000 0.000000 0.009375 0.000000 1.000000
##
##
## Variable -> Z_CostContact
```


Histograma de Z_CostContact




Boxplot de Z_CostContact




```
## [1] 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3
##   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##      3      3      3      3      3      3
##
##
## Variable ->  Z_Revenue
```

Histograma de Z_Revenue




Boxplot de Z_Revenue




```
## [1] 11 11 11 11 11 11 11 11 11 11 11 11 11 11 11 11 11 11
##   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##      11      11      11      11      11      11
##
##
## Variable ->  Response
```

Histograma de Response



Boxplot de Response



```
## [1] 1 0 0 0 0 0 0 0 1 0 0 0 0 0 0 1 0 0 0 0
##   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
## 0.0000 0.0000 0.0000 0.1491 0.0000 1.0000
```

Análisis de Variables Categóricas

```
# Seleccionar variables categóricas
categoricas <- sapply(ifood, function(x) is.factor(x) | is.character(x))
categoricas <- names(ifood)[categoricas]

# Análisis para cada variable categórica
for (var in categoricas) {
  cat("###", var, "\n\n")


  # Tabla de frecuencias
  print(table(ifood[[var]]))

  # Gráfico de barras
  barplot(table(ifood[[var]]), main=paste("Distribución de", var), col=rainbow(length(unique(ifood[[var]]))))


  # Gráfico de pastel
  pie(table(ifood[[var]]), main=paste("Distribución de", var), col=rainbow(length(unique(ifood[[var]]))))
}
```

```
## ### Education
##
##
##   2n Cycle   Basic Graduation   Master   PhD
##       203         54       1127       370       486
```

Distribución de Education



Distribución de Education



```
## ### Marital_Status
```


```
##
```

```
##
```


```
## Absurd Alone Divorced Married Single Together Widow YOLO
```

```
## 2 3 232 864 480 580 77 2
```

Distribución de Marital_Status



Distribución de Marital_Status



5. Preprocessing

Explain which preprocessing steps and methods you have applied to your dataset.

Load required libraries

```
# Suppress startup messages of library dplyr
suppressPackageStartupMessages(library(dplyr))
# Loading required libraries
library(dplyr, quietly = TRUE)
library(class, quietly = TRUE)
```

0. Load raw dataset

```
ifood <- read.csv("ml_project1_data.csv", sep=";", header=TRUE, stringsAsFactors = FALSE)
```

1. Remove irrelevant columns

```
ifood <- ifood[, !names(ifood) %in% c("ID", "Z_CostContact", "Z_Revenue")]
```

2. Transform date-related variables

```
ifood$Age <- 2020 - ifood$Year_Birth
ifood <- ifood[, !names(ifood) %in% c("Year_Birth")]
reference_date <- as.Date("2020-12-31")
ifood$CustDays <- as.numeric(reference_date - as.Date(ifood$Dt_Customer, format="%Y-%m-%d"))
ifood <- ifood[, !names(ifood) %in% c("Dt_Customer")]
```

3. Rename columns for easier access

```
colnames(ifood) <- gsub("NumDealsPurchases", "DealsPurc", colnames(ifood))
colnames(ifood) <- gsub("NumWebPurchases", "WebPurc", colnames(ifood))
colnames(ifood) <- gsub("NumStorePurchases", "StorePurc", colnames(ifood))
colnames(ifood) <- gsub("NumWebVisitsMonth", "WebVisits", colnames(ifood))
colnames(ifood) <- gsub("AcceptedCmpOverall", "CmpOverall", colnames(ifood))
colnames(ifood) <- gsub("MntWines", "WineExp", colnames(ifood))
colnames(ifood) <- gsub("MntFruits", "FruitExp", colnames(ifood))
colnames(ifood) <- gsub("MntMeatProducts", "MeatExp", colnames(ifood))
colnames(ifood) <- gsub("MntFishProducts", "FishExp", colnames(ifood))
colnames(ifood) <- gsub("MntSweetProducts", "SweetExp", colnames(ifood))
colnames(ifood) <- gsub("MntGoldProds", "GoldExp", colnames(ifood))
colnames(ifood) <- gsub("Marital_Status", "MaritalSts", colnames(ifood))
colnames(ifood) <- gsub("NumCatalogPurchases", "CatalogPurc", colnames(ifood))
colnames(ifood) <- gsub("AcceptedCmp1", "AccCmp1", colnames(ifood))
colnames(ifood) <- gsub("AcceptedCmp2", "AccCmp2", colnames(ifood))
colnames(ifood) <- gsub("AcceptedCmp3", "AccCmp3", colnames(ifood))
colnames(ifood) <- gsub("AcceptedCmp4", "AccCmp4", colnames(ifood))
colnames(ifood) <- gsub("AcceptedCmp5", "AccCmp5", colnames(ifood))
```

4. Handle outliers

```
ifood$Age <- ifelse(ifood$Age > 80, 80, ifood$Age)
```

5. Handle missing values

```
ifood <- ifood[!ifood$MaritalSts %in% c("YOLO", "Absurd"),]  
ifood$MaritalSts[ifood$MaritalSts == "Alone"] <- "Single"
```

6. Impute missing Income using KNN

```
ifood$Income <- ifelse(ifood$Income < 12500, NA, ifood$Income)  
  
num_vars <- sapply(ifood, is.numeric)  
complete_vars <- colnames(ifood)[num_vars]  
missing_threshold <- 0.2 * nrow(ifood)  
complete_vars <- complete_vars[colSums(is.na(ifood[, complete_vars])) < missing_threshold]  
aux <- ifood[, complete_vars]  
  
var <- "Income"  
aux1 <- aux[!is.na(ifood[[var]]), , drop = FALSE]  
aux2 <- aux[is.na(ifood[[var]]), , drop = FALSE]  
  
cols_na <- colnames(aux2)[colSums(is.na(aux2)) > 0]  
if (length(cols_na) > 0) {  
  aux1 <- aux1[, !(colnames(aux1) %in% cols_na), drop = FALSE]  
  aux2 <- aux2[, !(colnames(aux2) %in% cols_na), drop = FALSE]  
}  
  
knn_impute <- knn(aux1, aux2, ifood[[var]][!is.na(ifood[[var]])], k = 1)  
ifood[[var]][is.na(ifood[[var]])] <- as.numeric(as.character(knn_impute))
```

7. Correct calculation of TotAccCmp

```
ifood$TotAccCmp <- ifood$AccCmp1 + ifood$AccCmp2 + ifood$AccCmp3 + ifood$AccCmp4 + ifood$AccCmp5
```

8. Remove duplicate records

```
ifood <- ifood %>% arrange(desc(Response)) %>% distinct_at(vars(-Response), .keep_all = TRUE)
```

9. Create TotalExp before using it

```
ifood$TotalExp <- rowSums(ifood[, c("WineExp", "FruitExp", "MeatExp", "FishExp", "SweetExp", "GoldExp")], na.rm = TRUE)
```

10. Save cleaned dataset

```
write.csv(ifood, "ifood_cleaned.csv", row.names = FALSE)
```

Variable Creation

Second-Generation

Total Purchases

```
ifood$TotalPurchases <- ifood$DealsPurc + ifood$WebPurc + ifood$CatalogPurc + ifood$StorePurc
```

Purchase Frequency

```
ifood$PurchaseFrequency <- ifelse(ifood$CustDays > 0, ifood$TotalPurchases / (ifood$CustDays / 30), 0)
```

Preferred Product Category

```
product_categories <- c("WineExp", "FruitExp", "MeatExp", "FishExp", "SweetExp", "GoldExp")
max_index <- apply(ifood[, product_categories], 1, which.max)
ifood$PreferredProductCategory <- product_categories[max_index]
ifood$PreferredProductCategory <- as.factor(ifood$PreferredProductCategory)
```

Preferred Purchase Channel

```
channels <- c("DealsPurc", "WebPurc", "CatalogPurc", "StorePurc")
max_ch_index <- apply(ifood[, channels], 1, which.max)
ifood$PreferredChannel <- channels[max_ch_index]
ifood$PreferredChannel <- as.factor(ifood$PreferredChannel)
```

Average Spend Per Purchase

```
ifood$AvgSpendPerPurchase <- ifelse(ifood$TotalPurchases > 0, ifood$TotalExp / ifood$TotalPurchases, 0)
```

HasChildren

```
ifood$HasChildren <- ifelse(ifood$Kidhome + ifood$Teenhome > 0, 1, 0)
```

IncomeSegment

```
income_quantiles <- quantile(ifood$Income, probs = c(0.33, 0.66), na.rm = TRUE)
ifood$IncomeSegment <- cut(ifood$Income, breaks = c(-Inf, income_quantiles[1], income_quantiles[2], Inf),
  labels = c("Low", "Medium", "High"))
```

CustomerTenure

```
ifood$CustomerTenure <- ifood$CustDays / 365
```

CampaignAcceptanceRate

```
ifood$CampaignAcceptanceRate <- ifelse(ifood$TotAccCmp > 0, ifood$TotAccCmp / 5, 0)
```

Third-Generation

Third-Generation Feature 1: Customer Segmentation via Clustering

Prepare data for clustering: use Recency, TotalPurchases (frequency), and TotalExp (monetary)

```
cluster_data <- ifood %>% select(Recency, TotalPurchases, TotalExp)
```

Scale the data for clustering

```
cluster_data_scaled <- scale(cluster_data)
```

Perform k-means clustering with 3 clusters (as an example)

```
set.seed(123) # for reproducibility
k3 <- kmeans(cluster_data_scaled, centers = 3, nstart = 25) # nstart for better convergence
```

Add the cluster assignment as a new feature

```
ifood$CustomerSegment <- as.factor(k3$cluster)
```

(Customers are now labeled 1, 2, or 3 based on their cluster segment)

Third-Generation Feature 2: Propensity Score via Logistic Regression

Fit a logistic regression model to predict campaign response (Response) using relevant features

```
propensity_model <- glm(Response ~ Income + Recency + TotalExp + TotalPurchases + TotAccCmp + Age + MaritalSts,
  data = ifood, family = binomial)
```

Get predicted probabilities (propensity to respond)

```
ifood$PropensityScore <- predict(propensity_model, ifood, type = "response")
```

(PropensityScore is the model's predicted probability of Response=1 for each customer)

Quick summary of PropensityScore range

```
summary(ifood$PropensityScore)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
## 0.00000 0.03488 0.07725 0.15313 0.18046 0.99312
```

Third-Generation Feature 3: Engagement Index

Normalize components between 0 and 1

Note: For Recency, a lower value means more recent (more engaged), so we invert it.

```
recency_norm <- (max(ifood$Recency) - ifood$Recency) / max(ifood$Recency) # invert recency
frequency_norm <- ifood$TotalPurchases / max(ifood$TotalPurchases) # purchases normalized
monetary_norm <- ifood$TotalExp / max(ifood$TotalExp) # spending normalized
campaign_norm <- (ifood$TotAccCmp + ifood$Response) / 6 # campaign acceptance (out of 6 campaigns total including last response)
webvisit_norm <- ifood$WebVisits / max(ifood$WebVisits) # web visits normalized
```

Calculate engagement index as average of all five components, scaled to 0-100

```
ifood$EngagementIndex <- (recency_norm + frequency_norm + monetary_norm + campaign_norm + webvisit_norm) / 5 * 100
```

Preview EngagementIndex distribution

```
summary(ifood$EngagementIndex)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##  2.892  20.859  27.401  28.453  35.209  62.573
```

Save enriched dataset

```
write.csv(ifood, "ifood_enriched.csv", row.names = FALSE)
```

5.1 Explanation

The dataset preprocessing has been designed to ensure data quality, eliminate inconsistencies, and improve usability for analysis and modeling. Below, we explain the issues detected in the raw dataset and the reasoning behind each preprocessing step.

5.1.1 Issues Detected in the Raw Dataset

Before processing the data, we identified several issues that could affect the quality of the analysis:

1 Irrelevant Variables

- Some columns (**ID**, **Z_CostContact**, **Z_Revenue**) did not provide useful information for analysis.
- These columns did not contain relevant values for customer segmentation or behavior prediction.

2 Inadequate Formatting of Temporal Variables

- **Year_Birth**: Represented the customer's birth year. It was not directly interpretable, so it was transformed into **Age** to facilitate its use in models.

- **Dt_Customer**: Date of the first purchase in date format, which made certain analyses difficult. It was transformed into **CustDays** (days since the first purchase until 12/31/2020), converting it into a numerical metric for customer seniority.

3 Column Name Errors

- Some variable names were too long or confusing, making them difficult to use in code.
- They were renamed to more manageable versions. Example:
 - **NumWebPurchases** → **WebPurc**

4 Outlier Values

- **Age**: Extremely high values (>120 years) were found, indicating data entry errors or fictitious customers.
 - **Age** was capped at **80 years**, as the life expectancy in Brazil is approximately **76 years**, but many people exceed 80 (for example, women's life expectancy is 79 years).
 - This decision affected only **3 records**, preventing biases from erroneous data while keeping real elderly customers.
- **Income**: Abnormally low incomes (<12500, the annual minimum wage in Brazil in 2020) were detected, likely representing erroneous or outlier values in the distribution.

5 Missing Data and Invalid Values

- **MaritalSts**: Contained invalid values ("**YOLO**", "**Absurd**"), which did not represent legitimate marital status categories.
- **Income**:
 - Some values were **missing (NA)**.
 - Extremely low values (<12500) were considered erroneous and treated as missing data.

6 Calculation of TotAccCmp

- **TotAccCmp** was not correctly calculated. Instead of reflecting the total number of accepted campaigns, its value was incorrect.

7 Duplicate Records

- **Exact duplicate records** were found, which could bias the analysis if certain customers appear more times than they should.

5.1.2 Step-by-Step Preprocessing Explained

To address these issues, a structured pipeline for data cleaning and transformation was implemented.

1 Removal of Irrelevant Columns

JavaScript

```
ifood<-ifood[, !names(ifood) %in% c("ID",  
"Z_CostContact", "Z_Revenue")]
```

Reason: These variables do not provide useful information for the analysis.

- **ID** is a simple identifier.
- **Z_CostContact** and **Z_Revenue** appear to be constant or irrelevant.

2 Transformation of Temporal Variables

JavaScript

```
ifood$Age <- 2020 - ifood$Year_Birth  
ifood$CustDays<-as.numeric(reference_date-as.Date(ifood$Dt_Customer, format="%Y-%m-%d"))
```

Reason:

- **Year_Birth** is converted into **Age** because models better interpret age rather than birth year.
- **Dt_Customer** is transformed into **CustDays** (days since the first purchase) so that customer tenure becomes numerical and easy to use.

3 Renaming Columns for Clarity

JavaScript

```
colnames(ifood)<-gsub("NumWebPurchases", "WebPurc", colnames(ifood))
```

Reason: Some column names were too long (**NumWebPurchases** → **WebPurc**), making dataset writing and analysis more difficult.

4 Handling Outliers

JavaScript

```
ifood$Age<-ifelse(ifood$Age > 80, 80, ifood$Age)  
ifood$Income<-ifelse(ifood$Income<12500, NA, ifood$Income)
```

Reason:

- Ages over **80** are considered errors or unrepresentative in Brazil, but still allow for the inclusion of real elderly customers.
- Incomes below **12,500** are considered incorrect and are imputed as missing values.

5 Removal of Invalid Data in MaritalSts

JavaScript

```
ifood<-ifood[!ifood$MaritalSts %in% c("YOLO",  
"Absurd"), ]
```

Reason: **YOLO** and **Absurd** are not legitimate marital status categories, so they are removed from the dataset (only **4 records** were affected).

6 Imputation of Missing Values in Income

JavaScript

```
knn_impute<-knn(aux1,aux2,ifood[[var]][!is.na(ifood[[v  
ar]])], k = 1)  
ifood[[var]][is.na(ifood[[var]])]<-as.numeric(as.chara  
cter(knn_impute))
```

Reason: Instead of simply removing missing values in **Income**, they are imputed using **KNN**, leveraging information from other variables.

7 Correct Calculation of TotAccCmp

JavaScript

```
ifood$TotAccCmp<-ifood$AccCmp1+ifood$AccCmp2+ifood$Acc  
Cmp3 + ifood$AccCmp4 + ifood$AccCmp5
```

Reason: The variable was not correctly calculated. It now correctly reflects the **total number of accepted campaigns**

8 Removal of Duplicate Records

JavaScript

```
ifood<-ifood%>%arrange(desc(Response))%>%distinct_at(v  
ars(-Response), .keep_all = TRUE)
```

Reason: Duplicate records are removed, keeping only **one unique customer per row**. If duplicates exist, priority is given to those who responded **positively** to campaigns.

9 Creation of TotalExp Before Its Use

JavaScript

```
ifood$TotalExp <- rowSums(ifood[, c("WineExp",  
  "FruitExp", "MeatExp", "FishExp", "SweetExp",  
  "GoldExp")], na.rm = TRUE)
```

Reason: This variable **did not previously exist**, but it was required for several subsequent calculations.

5.2 Conclusion

This preprocessing process has been essential for correcting errors, improving the coherence of the dataset, and generating new variables useful for marketing analysis. The actions performed are detailed below:

- **Elimination of irrelevant information:** Variables and records that did not contribute value to the analysis were discarded, ensuring a cleaner and more focused dataset.
- **Transformation of dates and column names:** Date formats were standardized, and column names were renamed to facilitate interpretation and ensure better understanding of the data.
- **Correction of errors in age and atypical income:** Incorrect ages were adjusted, and income values that were out of range were normalized, improving the quality of the data for subsequent analysis.
- **Elimination of invalid values (YOLO, Absurd):** Records with anomalous values, such as "YOLO" or "Absurd", which were clearly erroneous data, were removed.
- **Imputation of missing values in "Income" using KNN:** An imputation process was applied to replace missing values in the "Income" field, using the KNN (K-Nearest Neighbors) algorithm, which ensures accurate estimation based on nearby data.
- **Correct calculation of "TotAccCmp":** The formula for "TotAccCmp" was corrected to ensure that the values were consistent and accurate, improving the integrity of the dataset.
- **Elimination of duplicates:** Duplicate records were removed to avoid bias in the analysis and ensure that each observation in the dataset is unique.
- **Creation of the "TotalExp" variable:** A new variable, "TotalExp", was created to facilitate the analysis of customers' total expenditures, enabling more effective segmentation and study of consumption behavior.

With these adjustments, the data is now ready to be used in prediction models and advanced segmentation.

5.3 Creation of New Variables:

Once the iFood dataset has been cleaned and refined, it is crucial to generate new variables to enhance the quality of the analysis and better align with the needs of data mining models (an idea based on the paper by Karina Gisbert: *A Survey on Pre-processing Techniques: Relevant Issues in the Context of Environmental Data Mining*).

The new variables are derived from the original ones using various techniques we have implemented and represent our proposed comprehensive preprocessing for the dataset.

5.3.1 Second-Generation Variables:

These refer to all variables based on expert knowledge, derived from domain-specific understanding. In our case, the “experts” are ourselves, as we have researched and studied the field to create these variables, allowing us to represent the concepts used by industry professionals in their reasoning.

These variables are derived by combining or transforming original data to better capture certain customer behaviors or characteristics.

- **TotalPurchases:** The total number of purchases a customer made across all channels.

Calculated as the sum of *DealsPurc + WebPurc + CatalogPurc + StorePurc*.

Why? This single metric captures overall purchase volume, indicating how active the customer is in transactions.

- **PurchaseFrequency:** Purchase frequency per month.

Calculated as: $PurchaseFrequency = TotalPurchases / (CustDays / 30)$, where *CustDays* represents the total number of days since the customer’s first purchase.

Why? This metric quantifies how often a customer makes a purchase, normalized per month so high values indicate frequent buyers, while low values suggest occasional or inactive customers. (e.g., useful for identifying loyal, high-frequency customers vs. sporadic buyers, enabling better-targeted retention and promotional strategies).

- **AverageSpendPerPurchase:** Average expenditure per transaction.

Computed as $TotalExp / TotalPurchases$ (the total monetary spend divided by the total number of purchases).

Why? This reveals whether the customer tends to make large purchases in each visit or smaller, lower-value transactions. It’s useful for identifying big spenders vs. bargain shoppers.

- **PreferredProductCategory:** Favorite product category based on spend.

For each customer, identify which of the product categories (WineExp, FruitExp, MeatExp, FishExp, SweetExp, GoldExp) is highest.

Why? This variable describes the customer's primary interest, allowing more personalized marketing – e.g., wine lovers vs. meat lovers have different profiles.

- **PreferredChannel:** Preferred shopping channel.

Determine the channel through which the customer made the most purchases: compare WebPurc, CatalogPurc, StorePurc. The new variable is a category like “Web”, “Store”, “Catalog”, or “Deals” for the channel with the highest purchase count for that customer. (We are considering Deals as another channel for purchasing.)

Why? This shows where the customer is most comfortable shopping, informing channel-specific strategies — e.g., online-oriented vs. in-store shoppers.

- **HasChildren:** Customers with children.

Calculated as a binary flag: *1 if TotalChildren > 0, else 0.*

Why? This explicitly distinguishes customers with families from those without, which could be insightful for segmentation.

- **IncomeSegment:** Customer income segment.

Categorizes customers into Low, Medium, or High income based on terciles of the Income distribution.

Calculated by computing income quantiles (33rd and 66th percentiles) and assigning each customer to one of the three segments:

- *Low:* Income in the bottom third.
- *Medium:* Income in the middle third.
- *High:* Income in the top third.

Why? This segmentation allows for differentiated marketing strategies based on spending power. — e.g., High-income customers might be more interested in gold products, or Low-income customers in promotions.

- **CustomerTenure:** Customer tenure in years. Measures how long the customer has been with the company.

Calculated as: $CustomerTenure = CustDays / 365$, converting days since first purchase into years.

Why? Identifies new vs. loyal customers for tailored retention strategies. e.g., Long-tenure customers may deserve loyalty rewards, while newer customers might need onboarding incentives.

- **CampaignAcceptanceRate:** Campaign acceptance rate. The percentage of marketing campaigns the customer has accepted.

Computed as: $CampaignAcceptanceRate = TotAccCmp / 5$, where TotAccCmp is the total number of accepted campaigns (sum of AccCmp1 to AccCmp5).

Why? High values indicate customers highly responsive to promotions, ideal for frequent marketing engagement. e.g., Low values signal less reactive customers, requiring stronger incentives or different approaches.

5.3.2 Third-Generation Variables:

These variables have been developed through independent research on various techniques, and it is essential to understand that this is a proposed approach. We consider it an additional enhancement to the given task (even though it is technically necessary).

We refer to variables generated using more sophisticated data analysis techniques (e.g., Clustering: k-means) to synthesize multiple original variables into more compact and useful indicators.

- **CustomerSegment:** A segment label for each customer, determined via clustering on key behaviors (e.g., spending, frequency, recency).

We have performed k-means clustering to group similar customers and assign a segment number.

Why? Identifies high-value active customers vs. low-value customers. (e.g., Targeted marketing)

- **PropensityScore:** A predictive score indicating the customer's propensity to respond to marketing campaigns.

We trained a simple logistic regression model using existing data (e.g., previous campaign acceptances and customer attributes) to estimate the probability of a positive response, and use this probability as the propensity score.

Why? Helps prioritize high-propensity customers for marketing campaigns.

EngagementIndex: A composite index that quantifies overall customer engagement.

This index will combine multiple factors (recency of purchases, frequency of purchases, total spending, campaign responsiveness, and website visits) into a single score (e.g., scaled 0–100).

Why? Captures overall customer involvement across purchases, campaigns, and web interactions.

We compute each of these step by step:

1. Customer Segmentation via Clustering

We use k-means clustering on key features to segment customers. Here we used RFM-style features: Recency, Frequency, and Monetary value (Recency, TotalPurchases, and TotalExp respectively). Before clustering, we scale these features to ensure equal weight. We choose 3 clusters for simplicity (this can be adjusted in the near future). After clustering, we assign each customer a segment label (1, 2, 3):

- Third-Generation Feature 1: Customer Segmentation via Clustering -

To prepare the data for clustering, we selected the features: Recency, TotalPurchases (frequency), and TotalExp (monetary). The next step is to scale the data for clustering purposes:

Python

```
# Prepare data for clustering: use Recency,
TotalPurchases (frequency), and TotalExp (monetary)
cluster_data <- ifood %>% select(Recency,
TotalPurchases, TotalExp)
# Scale the data for clustering
cluster_data_scaled <- scale(cluster_data)
# Perform k-means clustering with 3 clusters (as an
example)
set.seed(123) # for reproducibility
k3 <- kmeans(cluster_data_scaled, centers = 3, nstart
= 25) # nstart for better convergence
# Add the cluster assignment as a new feature
ifood$CustomerSegment <- as.factor(k3$cluster)
# (Customers are now labeled 1, 2, or 3 based on their
cluster segment)
```

We used 3 clusters to define broad segments (e.g., Segment 1, Segment 2, Segment 3). Depending on the clustering outcome, these might correspond to profiles such as “High-Value Customers”, “Occasional Buyers”, etc., but for now, they are simply numerical labels. (We converted `CustomerSegment` to a factor for clarity.)

2. Propensity Score (Predicted Response Probability)

To estimate each customer’s propensity to respond to marketing campaigns, we fit a logistic regression model using the cleaned dataset. The model predicts the probability of a customer accepting the next campaign (response variable) based on their characteristics (e.g., income, recency, past campaign acceptances, etc.). We then extract the predicted probability as the `PropensityScore` for each customer:

Python

```
# --- Third-Generation Feature 2: Propensity Score via
Logistic Regression ---
# Fit a logistic regression model to predict campaign
response (Response) using relevant features
propensity_model <- glm(Response ~ Income + Recency +
TotalExp + TotalPurchases + TotAccCmp + Age +
MaritalSts, data = ifood, family = binomial)
# Get predicted probabilities (propensity to respond)
ifood$PropensityScore <- predict(propensity_model,
ifood, type = "response")
# (PropensityScore is the model's predicted
probability of Response=1 for each customer)
# Quick summary of PropensityScore range
summary(ifood$PropensityScore)
```

In this case, we included features such as Income, Recency, TotalExpenditure, TotalPurchases, total past accepted campaigns, Age, and Marital Status as predictors. The resulting **PropensityScore** is a number between 0 and 1 (closer to 1 means the model predicts a higher likelihood the customer will respond positively to a campaign). In the near future, we might refine feature selection or use more advanced models.

3. Engagement Index

The **EngagementIndex** is designed to summarize how actively engaged a customer is with the brand across various dimensions. We combine multiple behaviors into one score. Specifically, we include:

- **Recency** (how recently the customer purchased, with more recent = more engaged),
- **Frequency** (total purchases),
- **Monetary** (total spending),
- **Campaign responsiveness** (number of campaigns accepted),
- **Web visit frequency** (number of web visits).

We will normalize each component to a 0–1 scale, then take an average (and scale to 0–100 for convenience):

Python

```
# --- Third-Generation Feature 3: Engagement Index ---
# Normalize components between 0 and 1
```

```

# Note: For Recency, a lower value means more recent
# (more engaged), so we invert it.
recency_norm <- (max(ifood$Recency) - ifood$Recency) /
max(ifood$Recency) # invert recency
frequency_norm <- ifood$TotalPurchases /
max(ifood$TotalPurchases) # purchases normalized
monetary_norm <- ifood$TotalExp / max(ifood$TotalExp)
# spending normalized
campaign_norm <- (ifood$TotAccCmp + ifood$Response) /
6 # campaign acceptance
(out of 6 campaigns total including last response)
webvisit_norm <- ifood$WebVisits /
max(ifood$WebVisits) # web
visits normalized
# Calculate engagement index as average of all five
# components, scaled to 0-100
ifood$EngagementIndex <- (recency_norm +
frequency_norm + monetary_norm + campaign_norm +
webvisit_norm) / 5 * 100
# Preview EngagementIndex distribution
summary(ifood$EngagementIndex)

```

In the code above, we treated a customer as more engaged if they have recent purchases (low **Recency** -> high **recency_norm**), frequent purchases (high **frequency_norm**), high spending (high **monetary_norm**), multiple accepted campaigns (**campaign_norm** accounts for 5 previous campaigns plus the latest response), and frequent web visits (high **webvisit_norm**). The final **EngagementIndex** is an average of these factors on a 0–100 scale. This provides a single metric to compare overall engagement levels across customers.

Conclusion of Third-Generation Features:

This set of third-generation variables is entirely our own proposal, developed through independent research. While not originally required, we have designed and implemented these variables to enhance the dataset by applying advanced data analysis techniques, ensuring a deeper and more structured understanding of customer behavior.

6. Basic description of post processing in this variables with changes

After preprocessing and enriching the dataset, several transformations have been applied to certain variables.

In general terms, the impact of preprocessing on the dataset has resulted in a reduction from 2.240 instances to 2.031 and an expansion from 29 variables to 35.

Original dataset:

```
# Cargar dataset
ifood <- read.csv("ml_project1_data.csv", sep=";", header=TRUE)

# Ver estructura del dataset
str(ifood)

## 'data.frame':    2240 obs. of  29 variables:
```

Enriched dataset:

```
# Cargar dataset
ifood <- read.csv("ifood_enriched.csv", sep=";", header=TRUE)

# Ver estructura del dataset
str(ifood)

## 'data.frame':    2031 obs. of  35 variables:
```

6.1 Main changes in variables

Based on the comparison of the analyses, the main changes we did in the following variables are:

- Numerical variables:
 - Income, Kidhome, Teenhome, and Recency: outlier correction and handling of missing values.
 - Product expenditures (WineExp, FruitExp, MeatExp, FishExp, SweetExp, GoldExp): summarized into a new variable (TotalExp).
 - Number of purchases per channel (WebPurc, StorePurc, CatalogPurc, DealsPurc): summarized into a new variable (TotalPurchases).
 - Accepted campaigns (AcceptedCmp1, AcceptedCmp2, AcceptedCmp3, AcceptedCmp4, and AcceptedCmp5): summarized into a new variable (TotalAccCmp).

- PurchaseFrequency, CustomerSegment, PropensityScore, and EngagementIndex: new derived variables added to the enriched dataset.
- Other variables were simply renamed to make them shorter (for example, WebVisitsMonth to WebVisits).
- Year_Birth was transformed into Age.
- Categorical variables:
 - Education and Marital_Status: outlier correction and handling of missing values.
 - PreferredProductCategory and PreferredChannel: new derived variables added to the enriched dataset.

6.2 Visualization of changes

In the analysis documents (created from “ifood_analysis.Rmd”), we have included histograms and boxplots for numerical variables, where we can observe the following:

- Histograms: more homogeneous distributions after preprocessing.
- Boxplots: reduction of outliers, indicating effective data cleaning.

Additionally, bar charts and pie charts are included for categorical variables, showing, after preprocessing, a more uniform distribution of previously redundant or residual instances.



6.3 Statistics’ summary after processing

After preprocessing, the descriptive statistics reveal:

- Reduction of extreme values: variables such as “Income” exhibit less extreme variability:



Original dataset:

##	###	Income					
##	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.	NA's
##	1730	35303	51382	52247	68522	666666	24



Enriched dataset:

##	###	Income					
##	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.	
##	12571	35828	51563	52844	68656	666666	




- Uniformization of categorical variable values, for example, in “Marital_Status” (merging of categories like “Single”, “Alone”, etc.):

Original dataset:

```
## ### Marital_Status
##
##
## Absurd Alone Divorced Married Single Together Widow
YOLO
## 2 3 232 864 480 580 77
```

Distribución de Marital_Status




Distribución de Marital_Status



Enriched dataset:

```
## ### MaritalSts
##
##
## Divorced Married Single Together Widow
## 213 787 445 516 70
```

Distribución de MaritalSts



Distribución de MaritalSts




- More balanced distributions: fewer outliers can be observed in the enriched dataset, such as in the “Year_Birth” variable, which was transformed into “Age”:


Original dataset:

```
## ### Year_Birth
##
## Min. 1st Qu. Median Mean 3rd Qu. Max.
## 1893 1959 1970 1969 1977 1996
```

Histograma de Year_Birth




Boxplot de Year_Birth




Enriched dataset:

```
## ### Age
##
## Min. 1st Qu. Median Mean 3rd Qu. Max.
## 24.0 43.0 50.0 51.2 61.0 80.0
```

Histograma de Age



Boxplot de Age




- New significant variables: “PropensityScore” and “EngagementIndex” provide additional analytical insights:


Enriched dataset:

```
## ### PropensityScore
##   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
## 0.00000 0.03488 0.07725 0.15313 0.18046 0.99312
```

Histograma de PropensityScore




Boxplot de PropensityScore




```
## ### EngagementIndex
##   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
## 2.892 20.859 27.401 28.453 35.209 62.573
```

Histograma de EngagementIndex



Boxplot de EngagementIndex



6.4 Conclusions

The adjustments we made have allowed for:

- A more precise and refined dataset for analysis.
- The introduction of new variables such as “EngagementIndex” and “PropensityScore” to generate more meaningful insights.
- Improved reliability of descriptive statistics, thanks to the handling of outliers and missing values.